



# GPU ACCELERATED COMPUTING WITH PYTHON

Sayak Bhowmick | Solutions Architect | December 2019

# WHY PYTHON?

- ❑ Very popular in many fields
  - ❑ Most preferred language for Deep Learning and Data Science
  - ❑ Scientific coding
- ❑ Features:
  - ❑ High-level, interactive, and interpreted
  - ❑ Garbage collection (reclaim memory from deleted variables)
  - ❑ Dynamically typed
  - ❑ Good data types and structures (intrinsic complex data type and Boolean)
  - ❑ Lots of libraries (modules) available
  - ❑ Easy to extend

# WHY PYTHON?

## And GPUs?

- ❑ You don't have to learn CUDA!
- ❑ High-level scripting languages are in many ways' polar-opposite to GPUs
  - ❑ GPUs are highly parallel, designed for maximum throughput, and they offer a tremendous advance in performance
  - ❑ Scripting languages such as Python favor ease of use over computational speed and do not generally emphasize parallelism
- ❑ Python coding for GPUs is getting better
  - ❑ Island for misfit Python/GPU code:
    - ❑ Lots of unsupported code hanging around
    - ❑ Not very Pythonic
  - ❑ No standards -> some standards (RAPIDS)
  - ❑ Interoperability is KEY

# TL;DR - OVERALL PYTHON-GPU CODING RECOMMENDATIONS

- ❑ These apply to any language coding for GPU, but are especially important for Python
- ❑ Avoid data movement to/from GPU and CPU
  - ❑ Do as much as you can on the GPU
- ❑ “What happens on the GPU, stays on the GPU”
- ❑ Look for loops, arrays
- ❑ When “porting” to GPU, intermediate code can be slower than CPUs
  - ❑ Don’t be surprised, don’t give up (profile)

# AGENDA

- ❑ Generalities
  - ❑ Quick, high-level recommendations
- ❑ Python Coding for GPUs
  - ❑ Numba (JIT compiler)
  - ❑ CuPy
  - ❑ Custom Kernels
    - ❑ CuPy
    - ❑ PyCuda

# GENERALITIES

## Scientific/Engineering Applications

- ☐ Good: Python has a huge number of libraries
- ☐ Bad: Python has a huge number of libraries
- ☐ Many scientific/engineering/data science codes are built using:
  - ☐ Numpy
  - ☐ SciPy
  - ☐ Scikit-learn
  - ☐ Pandas
- ☐ Much focus has been on Numpy for GPUs

**NUMBA**

# DIRECTIVES AND THE JIT

## JIT Compiler

- ❑ Language “Directives” tell compiler about code and how to build for target architecture (descriptive)
  - ❑ OpenACC for Fortran and C/C++ (compiled languages)
- ❑ Python is designed to be an interpreted language
  - ❑ No compilation or static typing
  - ❑ “Interactive”
- ❑ JIT = “Just in Time” Compiler
  - ❑ Compiles or creates object code “on the fly”
  - ❑ Combines interactivity and compilation
- ❑ Allows for computationally intensive sections of code to be run on GPU
  - ❑ Doesn’t have to be a function in NVIDIA libraries



# NUMBA

## Introduction

- ❑ Numba is a just-in-time (JIT), type-specializing, function compiler for accelerating numerically-focused Python
  - ❑ Not every Python function can be compiled (subset of Python and Numba)
  - ❑ Look for functions with high arithmetic intensity (loops of computations)
- ❑ Typically enabled by applying a *decorator* to a Python function
- ❑ Numba runs inside the standard Python interpreter
  - ❑ Can compile for CPU or GPU

# NUMBA

## Introduction - 2

- ❑ Uses LLVM to compile Python functions (comes with Numba)
- ❑ The first time the function is called, the compiler creates a machine code implementation for float inputs
  - ❑ Saves the original python implementation as `.py_func`
  - ❑ Can test compiled function against original Python function
- ❑ Subsequent calls to function use machine code (much faster)
  - ❑ You can create compiled code prior to running
- ❑ Data Types (dtypes):
  - ❑ `bool_`, `int8`, `int16`, `int32`, `int64`, `uint8`, `uint16`, `uint32`, `uint64`, `float16`, `float32`, `float64`
  - ❑ `numpy.complex64`, `numpy.complex128`

# NUMBA EXAMPLE

```
import numba
import math
```

```
@cuda.jit
```

```
def hypot(x, y):
```

```
    # From https://en.wikipedia.org/wiki/Hypot
```

```
    x = abs(x);
```

```
    y = abs(y);
```

```
    t = min(x, y);
```

```
    x = max(x, y);
```

```
    t = t / x;
```

```
    return x * math.sqrt(1+t*t)
```

Function Decorator

Function name/definition

Function to be compiled

# NUMBA EXAMPLE

```
import numpy as np
from numba import vectorize
```

```
@vectorize(['float32(float32, float32)'], target='cuda')
```

```
def Add(a, b):
    return a + b
```

```
N = 100000
```

```
# Initialize arrays
```

```
A = np.ones(N, dtype=np.float32)
```

```
B = np.ones(A.shape, dtype=A.dtype)
```

```
C = np.empty_like(A, dtype=A.dtype)
```

```
# Add arrays on GPU
```

```
C = Add(A, B)
```

@vectorize turns a scalar function to an elementwise array functions

Support multiple targets: 'cpu', 'parallel', 'gpu'

List of function type signatures

A scalar function

Numba takes care of data movement to/from CPU/GPU

# EXPLANATION

## What happened?

- ❑ Compiled a CUDA kernel to execute the `ufunc` operation in parallel over all the input elements
- ❑ Allocated GPU memory for the input(s) and the output
- ❑ Copied the input data to the GPU
- ❑ Executed the CUDA kernel with the correct kernel dimensions given the input sizes
- ❑ Copied the result back from the GPU to the CPU
- ❑ Returned the result as a NumPy object on the host

# ALLOWED NUMBA FUNCTIONS

(WARNING: It's not everything)

- ❑ Allowed statements/functions:

- ❑ `if/elif/else`
- ❑ `while` and `for` loops
- ❑ Basic math operators
- ❑ Selected functions from the `math` and `cmath` modules
- ❑ Tuples

<http://numba.pydata.org/numba-doc/latest/cuda/cudapysupported.html>

# NUMBA EXAMPLE 2

## Matrix Operations

- ❑ Compare Numpy, to Numba on CPU, to Numba on the GPU
- ❑ Compute  $\sin(x)\cos(x)$ 
  - ❑ Effectively loop over 1,000,000 values of  $x$
  - ❑ Numpy: Vector notation
  - ❑ Numba+CPU: vectorize function
  - ❑ Numba+GPU: vectorize function
    - ❑ Includes H -> D and D -> H

# NUMBA EXAMPLE 2

## Define Numba functions

```
# Get all the imports we need
import numba
import numpy as np
import math

# CPU version
@numba.vectorize(['float32(float32, float32)',
                 'float64(float64, float64)'], target='cpu')
def cpu_sincos(x, y):
    return math.sin(x) * math.cos(y)

# CUDA version
@numba.vectorize(['float32(float32, float32)',
                 'float64(float64, float64)'], target='cuda')
def gpu_sincos(x, y):
    return math.sin(x) * math.cos(y)
```

- ❑ Each function has two prototypes:
  - ❑ Float32
  - ❑ Float64



# NUMBA EXAMPLE 2

## Generate input data

```
# Generate data
n = 1000000
x = np.linspace(0, np.pi, n)
y = np.linspace(0, np.pi, n)

# Check result
np_ans = np.sin(x) * np.cos(y)
nb_cpu_ans = cpu_sincos(x, y)
nb_gpu_ans = gpu_sincos(x, y)

print("CPU vectorize correct: ", np.allclose(nb_cpu_ans, np_ans))
print("GPU vectorize correct: ", np.allclose(nb_gpu_ans, np_ans))
```

- ☐ Data is created on Host (CPU)
- ☐ Answers are generated on host as well (creates compiled functions)
- ☐ Answers are checked (functions work as expected)

# NUMBA EXAMPLE 2

Run benchmark (timeit)

```
print("NumPy")
%timeit np.sin(x) * np.cos(y)
```

```
print("CPU vectorize")
%timeit cpu_sincos(x, y)
```

```
print("GPU vectorize")
%timeit gpu_sincos(x, y)
```

```
# Optional cleanup
del x, y
```

## **NumPy**

10 loops, best of 3: 32 ms per loop

## **CPU vectorize**

10 loops, best of 3: 26.4 ms per loop

## **GPU vectorize**

10 loops, best of 3: 15.1 ms per loop

# JIT DECORATORS FOR GPU

- ❑ What is the difference between `@cuda.jit` and `@vectorize(target='gpu')`?
  - ❑ `@vectorize` will create ufuncs to take what were scalar inputs and allow vectors to be used
  - ❑ `@cuda.jit` creates PTX code and compiles it. It has CUDA related functions and variables.
- ❑ `@vectorize` is good for CPUs or learning how to port code to GPUs
- ❑ `@cuda.jit` is the best way to write functions that can use CUDA functions to tune compiling

CUPY

# CUPY

<https://github.com/cupy/cupy>

- ❑ NumPy-like API accelerated with CUDA
  - ❑ Implements a subset of NumPy, but it's very close to being complete
- ❑ Used by Chainer (DL framework popular in Japan)
- ❑ Calling sequence is like NumPy

```
>> import numpy as np
>> import cupy as cp
>> x_cpu = np.array([1, 2, 3])
>> l2_cpu = np.linalg.norm(x_cpu)
>> x_gpu = cp.array([1, 2, 3])
>> l2_gpu = cp.linalg.norm(x_gpu)
```
- ❑ Very Pythonic and very easy to use!

# CUPY

## Installation Notes

- ❑ Only available for Linux
- ❑ “`pip install cupy`”
  - ❑ Very old version
- ❑ “`conda install cupy`”

# CUPY FEATURES

- ❑ Patterned after Numpy
  - ❑ Doesn't include all Numpy functions, but a very high percentage
- ❑ Data Types (dtypes):
  - ❑ bool\_, int8, int16, int32, int64, uint8, uint16, uint32, uint64, float16, float32, float64
  - ❑ numpy.complex64, numpy.complex128
    - ❑ Test these since not every function may have it
  - ❑ Allows you to easily experiment with precision
- ❑ Can do custom kernels
- ❑ Documentation is pretty good

# CUPY EXAMPLE - 1

## Matrix Multiplication on GPU

```
import math
import cupy as cp

A = cp.random.uniform(low=-1., high=1., size=(64,64)).astype(cp.float32)
B = cp.random.uniform(low=-1., high=1., size=(64,64)).astype(cp.float32)

C = cp.matmul(A,B)
```

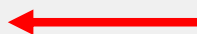


# CUPY EXAMPLE - 2

## SVD

```
import cupy as cp  
  
A = cp.random.uniform(low=-1., high=1., size=(64, 64)).astype(cp.float32)  
u, s, v = cp.linalg.svd(A)
```

u, s, v are  
still on GPU



# CUPY EXAMPLE - 3

## SVD - 2 (copy data back to CPU)

```
import cupy as cp
import numpy as np
```

```
A_cpu = np.random.uniform(low=-1., high=1., size=(64, 64)).astype(np.float32)
A_gpu = cp.asarray(A_cpu)
```

Copy A\_cpu to GPU.  
Becomes CuPy object

```
u_gpu, s_gpu, v_gpu = cp.linalg.svd(A_gpu)
print ("type(u_gpu) = ", type(u_gpu))
```

```
u_cpu = cp.asnumpy(u_gpu)
print ("type(u_cpu) = ", type(u_cpu))
```

Copy u\_gpu to Host.  
Becomes Numpy object

```
[sayakb@hsw225 cupy]$ python3 svd2.py
type(u_gpu) = <type 'cupy.core.core.ndarray'>
type(u_cpu) = <type 'numpy.ndarray'>
```

# CUSTOM KERNELS

# CUSTOM KERNELS

- ❑ What do you do if your needs are not met by any of the libraries/tools?
- ❑ Perhaps it's time to learn `_some_` CUDA and write your own custom kernel
  - ❑ CuPy
  - ❑ PyCUDA

# CUPY CUSTOM KERNELS

- ❑ Four types of kernels:
  - ❑ `cupy.ElementwiseKernel`
  - ❑ `cupy.ReductionKernel`
  - ❑ `cupy.RawKernel`
  - ❑ `cupy.fuse`
- ❑ Most custom kernels come with some (many?) CUDA functions

# CUPY ELEMENTWISE KERNEL

- ❑ Focuses on kernels that operate on an element-wise basis
- ❑ Consist of 4 components:
  - ❑ input argument list (comma-separated)
  - ❑ output argument list (comma-separated)
  - ❑ loop body code
  - ❑ kernel name

```
import cupy as cp

kernel = cp.ElementwiseKernel(
    'float32 x, float32 y', 'float32 z',
    '''if (x - 2 > y) {
        z = x * y;
    } else {
        z = x + y;
    }''', 'my_kernel')
```

# CUPY REDUCTION KERNEL

- ❑ Has four parts:
  - ❑ Identity value: Initial value of the reduction
  - ❑ Mapping expression: Preprocesses each element to be reduced
  - ❑ Reduction expression: An operator to reduce the multiple mapped values. Two special variables, `a` and `b`, are used for this operand
  - ❑ Post-mapping expression: Transforms the reduced values. The special variable `a` is used as input. The output should be written to the output variable.
- ❑ This function uses type placeholders for both the input and output

```
import cupy as cp

l2norm_kernel = cp.ReductionKernel(
    'T x', # input params
    'T y', # output params
    'x * x', # map
    'a + b', # reduce
    'y = sqrt(a)', # post-reduction map
    '0', # identity value
    'l2norm' # kernel name
)

x = cp.arange(10, dtype=np.float32).reshape(2, 5)
l2norm_kernel(x, axis=1)

array([ 5.477226 , 15.9687195], dtype=float32)
```

# CUPY RAW KERNEL

- ❑ Define using “raw” CUDA code
- ❑ Lacks helpful variables or CuPy functions or elementwise or reduction custom functions

```
import cupy as cp

add_kernel = cp.RawKernel(r'''
    extern "C" __global__
    void my_add(const float* x1, const float* x2, float* y) {
        int tid = blockDim.x * blockIdx.x + threadIdx.x;
        y[tid] = x1[tid] + x2[tid];
    }
    ''', 'my_add')

x1 = cp.arange(25, dtype=cp.float32).reshape(5, 5)
x2 = cp.arange(25, dtype=cp.float32).reshape(5, 5)
y = cp.zeros((5, 5), dtype=cp.float32)
add_kernel((5,), (5,), (x1, x2, y)) # grid, block and arguments

print(y)

array([[ 0.,  2.,  4.,  6.,  8.],
       [10., 12., 14., 16., 18.],
       [20., 22., 24., 26., 28.],
       [30., 32., 34., 36., 38.],
       [40., 42., 44., 46., 48.]], dtype=float32)
```



# CUPY FUSE

- ❑ This decorator can be used to define an elementwise or reduction kernel more easily than `ElementwiseKernel` or `ReductionKernel`

```
@cupy.fuse()
def squared_diff(x, y):
    return (x - y) * (x - y)

x = cupy.arange(10)
y = cupy.arange(10)[::-1]

squared_diff(x, y)

array([81, 49, 25,  9,  1,  1,  9, 25,
       49, 81])
```

# PYCUDA - CUSTOM KERNELS

- ❑ Earliest integration of Python and CUDA but requires a knowledge of CUDA and Python
- ❑ Really a tool to compile and interface CUDA code with Python
  - ❑ It does have some Python variables that correspond to CUDA variables
  - ❑ Functions to make GPU coding easy (<https://document.tician.de/pycuda/util.html>)
    - ❑ Example: Copying data to/from GPU
- ❑ PyCUDA's base layer is written in C++

# PYCUDA SIMPLE EXAMPLE

- ❑ Generate random numbers on CPU
- ❑ Copy to GPU
- ❑ Double the values in the array in parallel
- ❑ Copy the data back to CPU

# SIMPLE PYCUDA

- ❑ `pycuda.autoinit` is used for automatic initialization, context creation, and cleanup
- ❑ Numpy is used to generate a random 4x4 array
- ❑ Allocate space (`mem_alloc`) on GPU
- ❑ Copy CPU data to GPU
  - ❑ Use built-in PyCUDA function `memcpy_htod()`
- ❑ Define “C” code using `SourceModule`
- ❑ Compile function (calls `nvcc`)
- ❑ Call function with “execution configuration” (grid properties)
- ❑ Copy data back to CPU

```
import pycuda.driver as cuda
import pycuda.autoinit
from pycuda.compiler import SourceModule

import numpy as np

a = np.random.randn(4, 4)
a = a.astype(np.float32) # convert to FP32

a_gpu = cuda.mem_alloc(a.nbytes) # Malloc on GPU
cuda.memcpy_htod(a_gpu, a)      # Copy data to GPU

mod = SourceModule(
    """__global__ void doublify(float *a)
    {
        int idx = threadIdx.x + threadIdx.y * 4;
        a[idx] *= 2;
    }
    """)

func = mod.get_function("doublify") # compile function
func(a_gpu, block = (4, 4, 1))     # "call" function, 4x4 grid

a_doubled = np.empty_like(a) # Create data for results
cuda.memcpy_dtoh(a_doubled, a_gpu)
```

# SUMMARY

# SUMMARY

- ❑ Coding GPUs for Python can be very easy -or- more complex, but with more control
  - ❑ Compile Python code using Numba (almost 100% pure python) - don't need to know CUDA!
  - ❑ Use Numpy like functions on GPU with CuPy - you don't necessarily need to know CUDA!!
  - ❑ Custom kernels with CuPy and PyCUDA - need to know CUDA!!!
- ❑ Many of these tools are interoperable
  - ❑ The start of making GPU's first-class citizens with Python

# GET STARTED TODAY

You might already have a CUDA-capable GPU in your laptop or desktop PC!

NVIDIA Developer Zone

<https://developer.nvidia.com/>

CUPY

<https://cupy.chainer.org/>

NUMBA

<http://numba.pydata.org/numba-doc/latest/cuda/index.html>

PYCUDA

<https://documen.tician.de/pycuda/index.html>

**sayakb@nvidia.com**

The background is a dark blue gradient with a complex network of thin, glowing green lines. These lines connect various points, some of which are larger, bright green circular nodes. The overall effect is a sense of a dynamic, interconnected system or network.

**Q & A**



